Lab 1: Elastic Observability

In these labs, you will get access to a system, which we will refer to as "host". On this host, an Nginx web server is already running. You will learn how you can use Elastic Observability to monitor and gain insights into Nginx and this host.

In labs 2 and 3, you will learn how to deploy and configure Elastic Agent. Next, you will use the Nginx integration to collect logs and metrics from Nginx. You will collect logs and metrics from the host on which Nginx is running as well.

In lab 4 you will explore APM to collect detailed performance information.

Objective:

In this lab, you are going to familiarize yourself with the Elastic Cloud Elasticsearch Service. The Elastic Cloud Elasticsearch Service is an Elasticsearch-as-a-Service product which eases the creation, configuration, and maintenance of large and complex infrastructures.

1. First of all, you need to log in to Elastic Cloud Elasticsearch Service. Click here to open the main page. If you already have a cloud account just log in. If you do not have one, click **Sign up**, enter your email as well as a (strong!) password and click **Create account**.

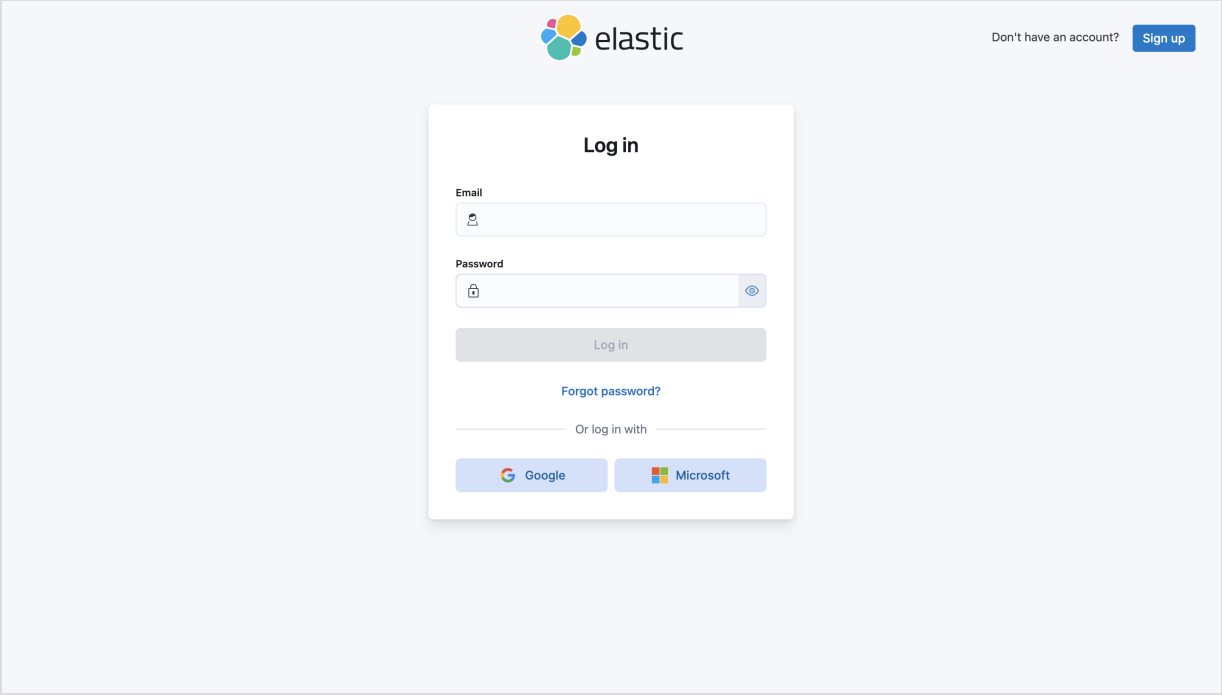

2. After logging in, you will reach the Elastic Cloud console. There you will find information about deployments, news, training, documentation, and webinars. Note that you do not have any deployments because you just created your account. Click on **Start your free trial** (or **Create deployment**) to create your first deployment.

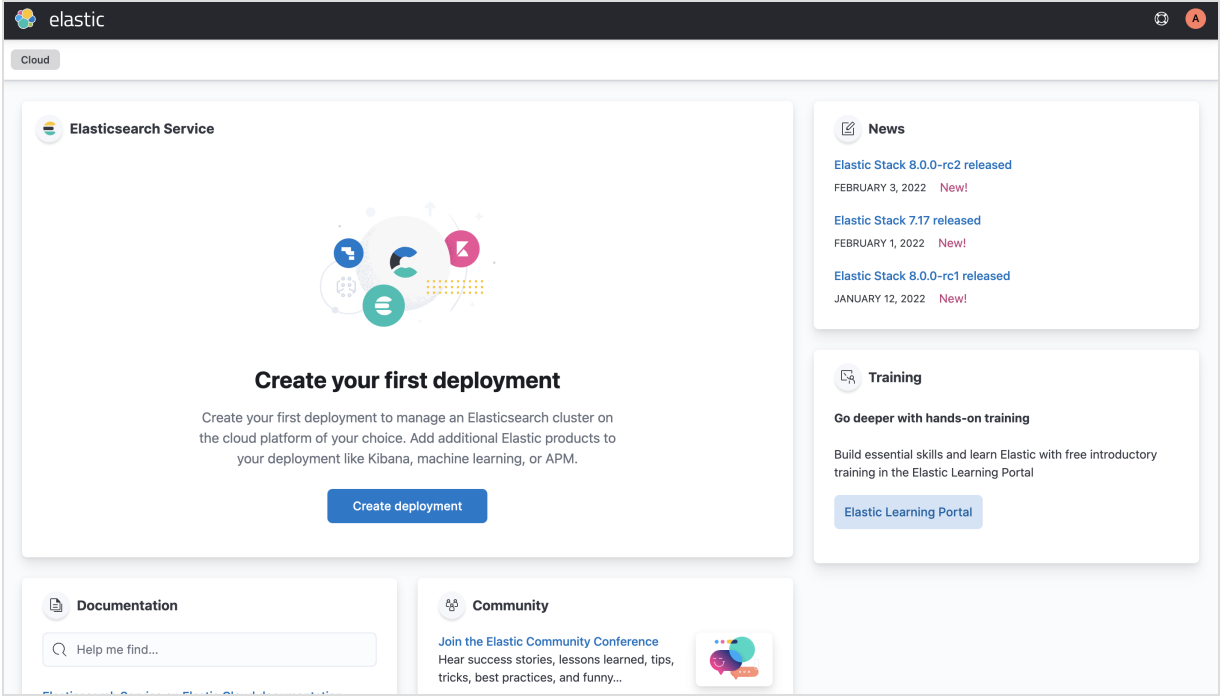

3. To create your first deployment you will need to give it a meaningful name. For this training we will be calling it `observability-deployment`.

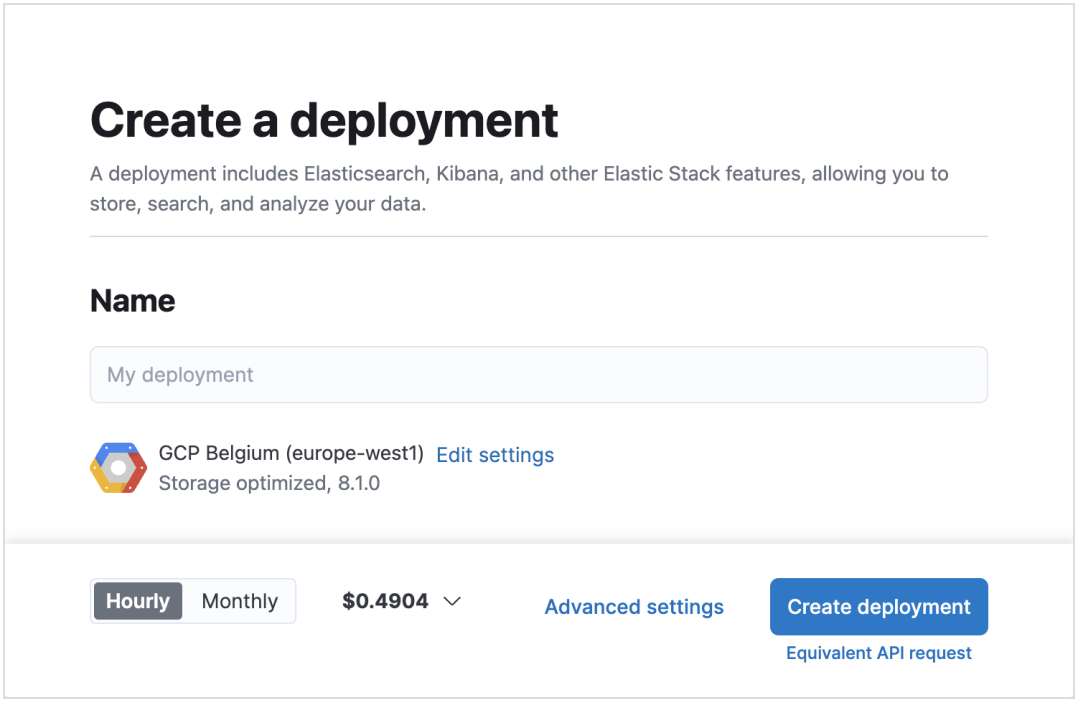

4. You will also need to select a cloud provider and region, choose a hardware profile, and define the version that you will be using. For this training select **Google Cloud**, but note that you might also select either AWS or Azure for your future deployments. Also select a region that is close to you and the latest version that is available. For hardware profile you can use **Storage Optimized**, which is ideal for observability data.

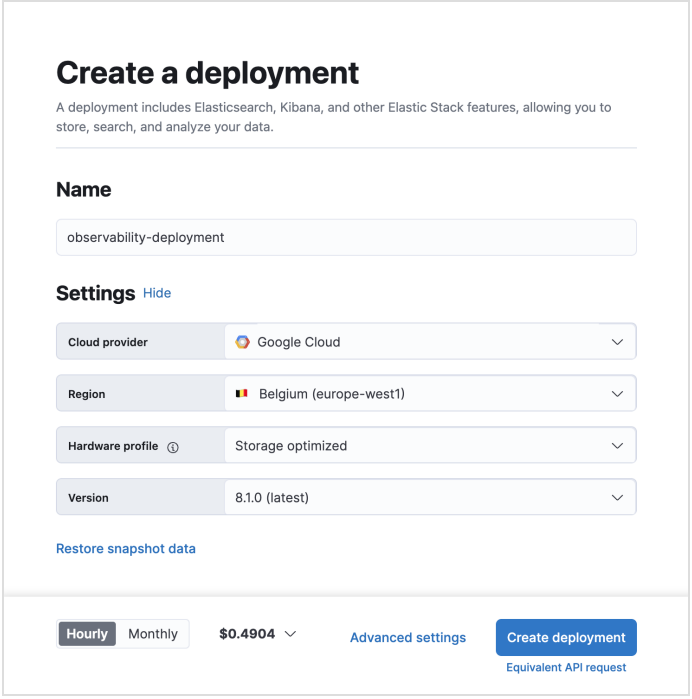

5. When you are all set click **Create deployment**.
6. Once you click **Create deployment**, you are prompted to save the admin credentials for your deployment. You should download or copy the credentials to a safe place, as they are shown only once.

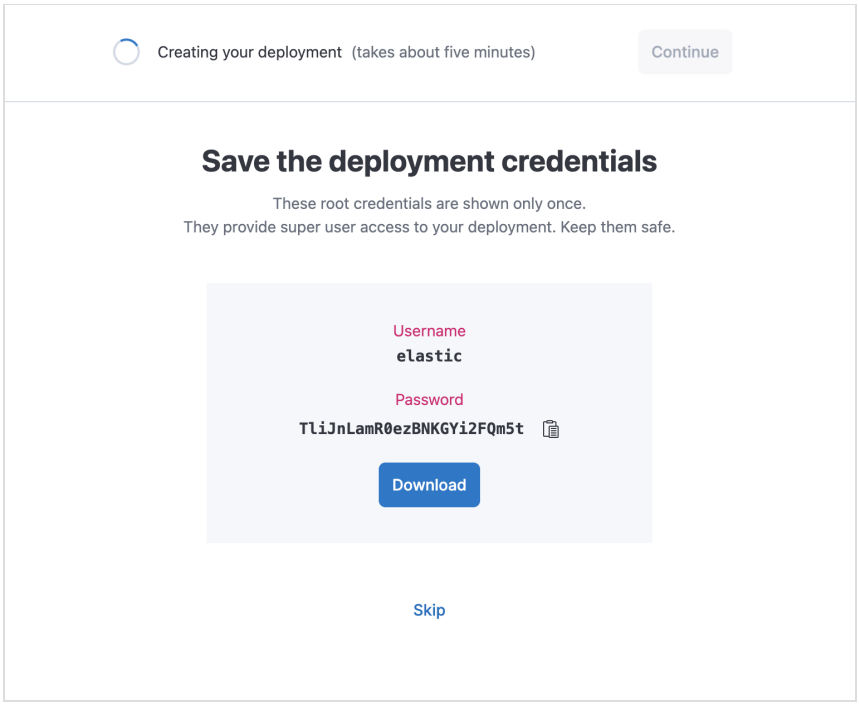

7. It will take a few minutes until your deployment is up and running. When it is ready click on **Continue** to open Kibana.
8. After successfully logging in to Kibana you should see a landing page, which looks like this:

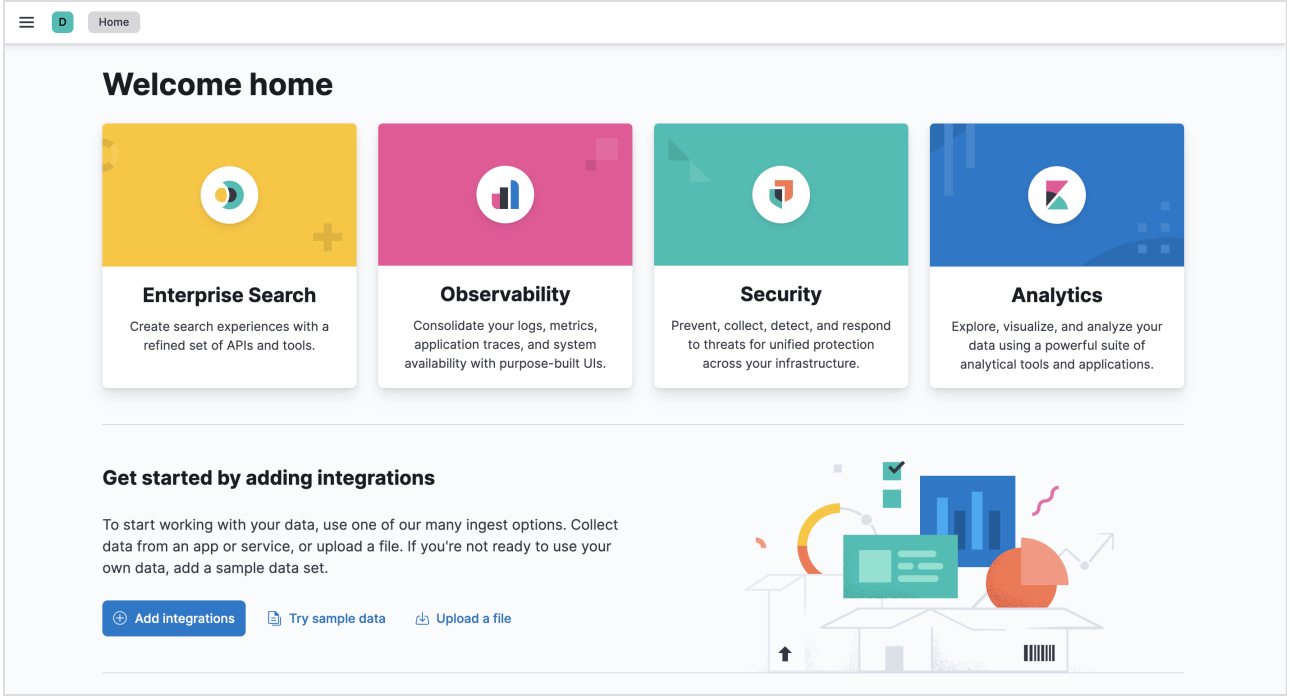

9. Feel free to explore Kibana.
10. Return to the Elastic Cloud console.
11. If you close Kibana and want to get back to it, you can click the deployment name from the Elastic Cloud console.
12. Click the gear shown under **Quick links** for your `observability-deployment`. Here, you will see some information about the deployment.

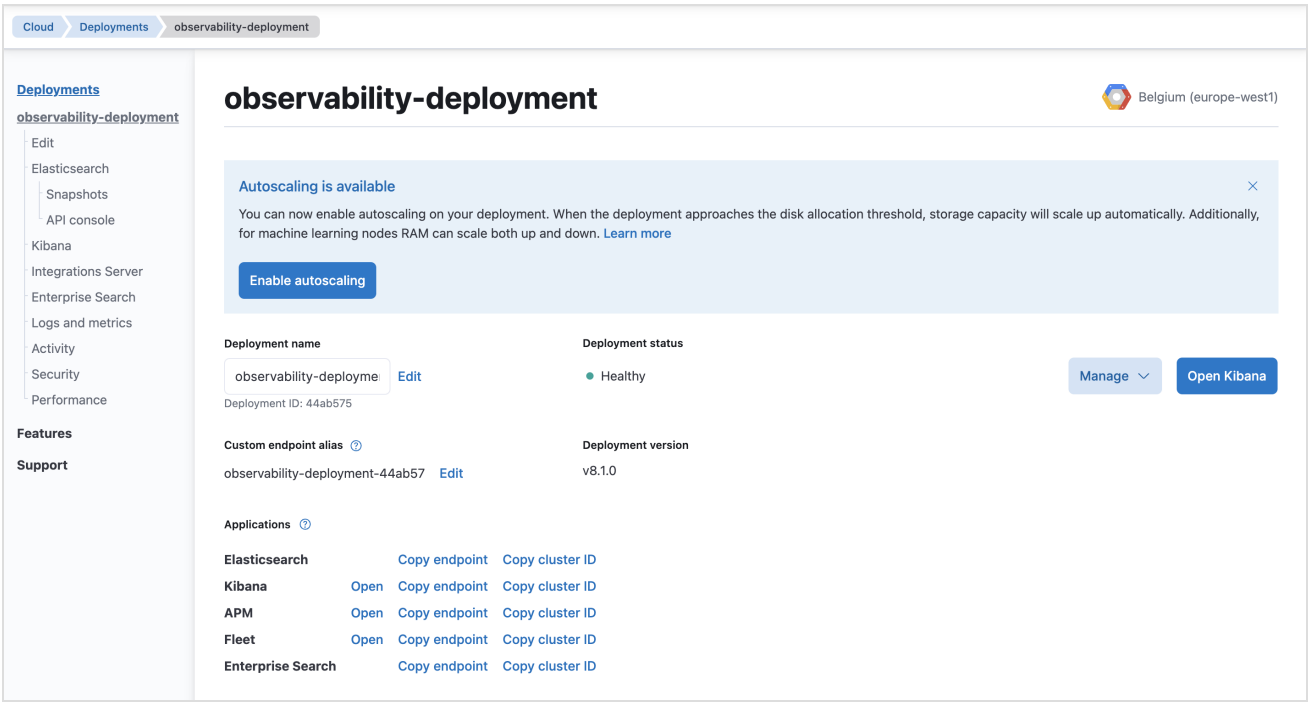

13. Scroll down to check the information about your deployment and answer some questions about it.

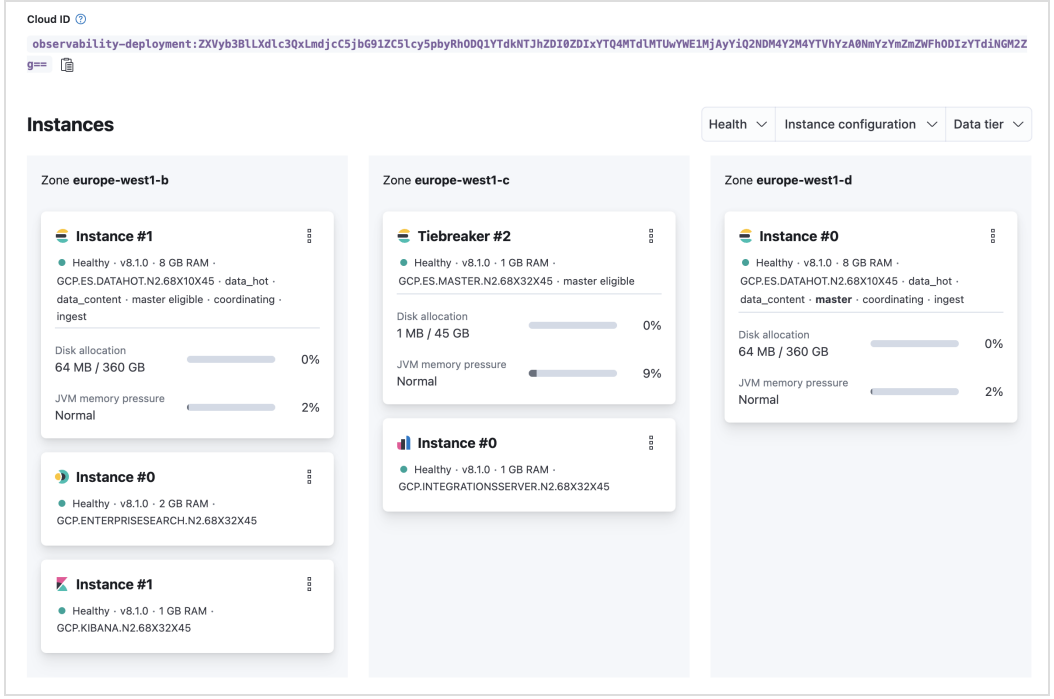

- **What is the Elasticsearch endpoint of your deployment?**

  Click on the *Copy endpoint* link next to the Elasticsearch label top copy the Elasticsearch endpoint. You might need it if you want to access Elasticsearch through its RESTful APIs.

- **What is the Kibana endpoint?**

  Click on the *Copy endpoint* link next to the Kibana label. You can use the *Open* link to open Kibana.

- **How many Elasticsearch nodes are in your cluster?**

  There are 3 Elasticsearch nodes in our example: `Instance #0`, `Instance #1`, and `Tiebreaker #2`.

  In addition to the information above, here you can also find the version of your cluster, how much memory each node has, the JVM memory pressure, the current disk usage, etc.

Summary:

You created an Elastic Cloud deployment in just a few clicks and now it's ready to start indexing observability data!